

Decision Trees

-Entropy = gain information

Tian-Li Yu

Taiwan Evolutionary Intelligence Laboratory (TEIL)
Department of Electrical Engineering
National Taiwan University
tianliyu@ntu.edu.tw

Readings: AIMA 18.3, 18.4, 18.10

Outline

- 1 Learning Decision Trees
 - Choosing Attributes
- 2 Model Evaluation
 - Metrics
 - Cross-Validation
 - Comparison
- 3 Generalization and Overfitting
- 4 Ensemble

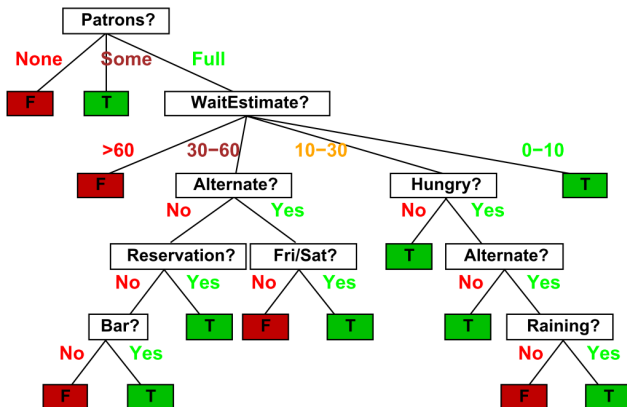
Attribute-based Representations

- Restaurant example.

Example	Attributes										Target
	<i>Alt</i>	<i>Bar</i>	<i>Fri</i>	<i>Hun</i>	<i>Pat</i>	<i>Price</i>	<i>Rain</i>	<i>Res</i>	<i>Type</i>	<i>Est</i>	<i>WillWait</i>
X_1	<i>T</i>	<i>F</i>	<i>F</i>	<i>T</i>	<i>Some</i>	<i>\$\$\$</i>	<i>F</i>	<i>T</i>	<i>French</i>	<i>0-10</i>	<i>T</i>
X_2	<i>T</i>	<i>F</i>	<i>F</i>	<i>T</i>	<i>Full</i>	<i>\$</i>	<i>F</i>	<i>F</i>	<i>Thai</i>	<i>30-60</i>	<i>F</i>
X_3	<i>F</i>	<i>T</i>	<i>F</i>	<i>F</i>	<i>Some</i>	<i>\$</i>	<i>F</i>	<i>F</i>	<i>Burger</i>	<i>0-10</i>	<i>T</i>
X_4	<i>T</i>	<i>F</i>	<i>T</i>	<i>T</i>	<i>Full</i>	<i>\$</i>	<i>F</i>	<i>F</i>	<i>Thai</i>	<i>10-30</i>	<i>T</i>
X_5	<i>T</i>	<i>F</i>	<i>T</i>	<i>F</i>	<i>Full</i>	<i>\$\$\$</i>	<i>F</i>	<i>T</i>	<i>French</i>	<i>>60</i>	<i>F</i>
X_6	<i>F</i>	<i>T</i>	<i>F</i>	<i>T</i>	<i>Some</i>	<i>\$\$</i>	<i>T</i>	<i>T</i>	<i>Italian</i>	<i>0-10</i>	<i>T</i>
X_7	<i>F</i>	<i>T</i>	<i>F</i>	<i>F</i>	<i>None</i>	<i>\$</i>	<i>T</i>	<i>F</i>	<i>Burger</i>	<i>0-10</i>	<i>F</i>
X_8	<i>F</i>	<i>F</i>	<i>F</i>	<i>T</i>	<i>Some</i>	<i>\$\$</i>	<i>T</i>	<i>T</i>	<i>Thai</i>	<i>0-10</i>	<i>T</i>
X_9	<i>F</i>	<i>T</i>	<i>T</i>	<i>F</i>	<i>Full</i>	<i>\$</i>	<i>T</i>	<i>F</i>	<i>Burger</i>	<i>>60</i>	<i>F</i>
X_{10}	<i>T</i>	<i>T</i>	<i>T</i>	<i>T</i>	<i>Full</i>	<i>\$\$\$</i>	<i>F</i>	<i>T</i>	<i>Italian</i>	<i>10-30</i>	<i>F</i>
X_{11}	<i>F</i>	<i>F</i>	<i>F</i>	<i>F</i>	<i>None</i>	<i>\$</i>	<i>F</i>	<i>F</i>	<i>Thai</i>	<i>0-10</i>	<i>F</i>
X_{12}	<i>T</i>	<i>T</i>	<i>T</i>	<i>T</i>	<i>Full</i>	<i>\$</i>	<i>F</i>	<i>F</i>	<i>Burger</i>	<i>30-60</i>	<i>T</i>

Decision Trees

- One possible representation for hypotheses.



Expressiveness of Decision Trees

- $Goal \Leftrightarrow (Path_1 \vee Path_2 \vee \dots)$.
- $Path_i \Leftrightarrow (Attribute_1 = a_1 \wedge Attribute_2 = a_2 \wedge \dots)$.
- Decision trees can express **any** function of the input attributes.
- Trivially, there is a consistent decision tree for **any** training set with one path to leaf for each example, but it won't generalize to new examples.
- Prefer to find more **compact** decision trees.

Hypothesis Space

- How many **distinct** decision trees with n Boolean attributes?
 - Truth table with 2^n rows.
 - Every truth table can be expressed by one decision tree \Rightarrow At least 2^{2^n} decision trees.
 - If different order of attributes counts as \Rightarrow At least $n! \cdot 2^{2^n}$ decision trees.
- More expressive hypothesis space
 - Increase the chance that c can be expressed.
 - May be weak at generalization if we let the decision tree be too expressive.

Learning Decision Trees (ID3 [Quinlan, 1986])

- Aim: Find a **small** tree **consistent** with training examples.
- Idea: Recursively choose the **best** attribute.

DTL(*examples*, *attributes*, *examples_{parent}*)



```

1  if examples is empty then return PLURALITY-VALUE(examplesparent)
2  elseif all examples have same classification then return the classification
3  elseif attributes is empty then return PLURALITY-VALUE(examples)
4  else
5       $A \leftarrow \operatorname{argmax}_{a \in \text{attributes}} \text{IMPORTACE}(a, \text{examples})$ 
6      tree  $\leftarrow$  a new decision tree with root A
7      for textbf each value  $v_k$  of A
8          exs  $\leftarrow$  elements of examples with  $A = v_k$ 
9          subtree  $\leftarrow$  DTL(exs, attributes – A, examples)
10         add a branch to tree with label  $A = v_k$  and subtree subtree
11     return tree

```

Choosing Attributes

- The restaurant example consist of 6 positive and 6 negative examples.
- *Patrons* is a better choice — gives more **information** about the classification.



Information

- Measure of information: **Shannon's entropy**.
 - Gives the lower bound of the most compact encoding of a random variable in bits.
- The entropy of a random variable V with values v_k , each with probability $P(v_k)$, is defined as

$$H(V) = - \sum_k P(v_k) \log_2 P(v_k).$$

- For Boolean variables, define

$$B(q) = -q \log_2 q - (1 - q) \log_2 (1 - q).$$

- The entropy of a fair coin:
 $H = B(0.5) = -(0.5 \log_2 0.5 + 0.5 \log_2 0.5) = 1 \text{ bit}.$
- The entropy of a unfair coin (99% head):
 $H = B(0.99) = -(0.99 \log_2 0.99 + 0.01 \log_2 0.01) = 0.08 \text{ bits}.$

Information

- p positive and n negative examples at the root $\Rightarrow B(p/(p+n))$ bits needed to classify a new example.
- Attribute A splits the examples E into subsets E_k , each of which (we hope) needs less information to complete the classification.
- Let E_k have p_k positive and n_k negative examples
 $\Rightarrow B(p_k/(p_k+n_k))$ bits needed to classify a new example
 \Rightarrow **expected** number of bits per example over all branches is

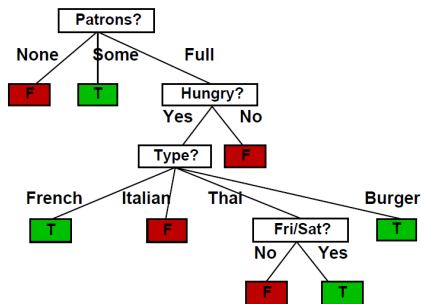
$$\text{Remainder}(A) = \sum_k \frac{p_k+n_k}{p+n} B\left(\frac{p_k}{p_k+n_k}\right).$$

- For *Patrons*, this is 0.459 bits; for *Type*, this is (still) 1 bit
 \Rightarrow Choose the attribute that minimizes the remaining information.
 \Rightarrow Choose the attribute with the most **information gain**:

$$\text{Gain}(A) = B\left(\frac{p}{p+n}\right) - \text{Remainder}(A)$$

Decision Tree Learned from the Examples

- Decision tree learned from the 12 examples:



- Substantially simpler than a full tree — a more complex hypothesis isn't justified by small amount of data.

Model Evaluation

- Metrics for Performance Evaluation
 - How to evaluate the performance of a model?
- Methods for Performance Evaluation
 - How to obtain reliable estimates?
- Methods for Model Comparison
 - How to compare the relative performance among competing models?

Metrics for Performance Evaluation

- Confusion matrix:

ACTUAL CLASS	PREDICTED CLASS		
		Class=Yes	Class=No
	Class=Yes	True Positive	False Negative
	Class=No	False Positive	True Negative

Accuracy

$$Accuracy = \frac{TP + TN}{TP + TN + FP + FN}$$

對角線最大化約好

- Probably most widely-used metric.
- Can be misleading. Consider class 0 consisting of 9990 instances and class 1 consisting of 10 instances. Classifying everything as class 0 yields 99.9% accuracy.

當even時有缺點

Other Metrics


$$\textit{Precision}(p) = \frac{TP}{TP + FP}$$

$$\textit{Recall}(r) = \frac{TP}{TP + FN}$$

$$F - \textit{measure}(F) = \frac{2pr}{p + r} = \frac{2TP}{2TP + FP + FN}$$

Performance Measurement

- How do we know whether $h \approx c$?
 - ① Use theorems of computational/statistical learning theory
 - ② Try h on a new **test set** of examples
(use **same distribution** over example space as training set)
- **Learning curve** = % correct on **test set** as a function of **training set size**



Cross-Validation

data 不會均分

test data < train data (e.g. 1:3, 1:4...)

- The idea of having training and testing sets is called **cross-validation**.
- **Holdout cross-validation**
 - **Randomly split** the available data into a training set and a testing set.
 - Simple, fast, but not able to use all available data.
- **k-fold cross-validation**

random均分成k份
第n次用第n組當testing data, 剩下k-1組training data

 - **Randomly split** the data into k equal-sized subsets.
 - Perform k rounds of learning using $k - 1$ subsets as training and the rest as testing.
 - Popular choice of k is 5 to 10.
 - Accurate statistics, but longer computation.

ROC

- Receiver operating characteristic.
- ROC curve: FPR as x -axis; TPR as y -axis.
false positive rate


$$FPR(FP \text{ rate}) = \frac{FP}{FP + TN}$$

$$TPR(TP \text{ rate}) = \frac{TP}{TP + FN}$$

- (FPR, TPR):
 - (0,0): Classify everything as negative.
 - (1,1): Classify everything as positive.
 - (0,1): Ideal. 沒有東西弄錯

AUC

- Model 1 is better for small FPR.
- Model 2 is better for large FPR.
- Area under the ROC curve (AUC).
 - Ideal: 1.
 - Random guess: 0.5.



Generalization and Overfitting

- If some attributes are irrelevant, DTL still outputs a large tree.
 - The outputs of fair dices with attributes of color, size, and so on.
- To overcome **overfitting**,
 - we can stop growing the tree before overfitting,
 - or we can allow overfitting, and then post-prune the tree (most common).
- How to decide what to post-prune?
 - Use the testing set. training 長樹
砍了某個以後testing有變好 就繼續post-prune
 - Use statistical tests.
 - Use explicit measures the complexity of the encoding of the tree and training examples (minimum description length principle).

χ^2 Pruning

- **Information gain** of an irrelevant attribute is **expected** to be zero, but the **sampling noise** may still yield some gain.
- Assuming true irrelevant, the expected number of p_k and n_k can be expressed as

$$\hat{p}_k = \frac{p}{p+n} \times (p_k + n_k) \quad \hat{n}_k = \frac{n}{p+n} \times (p_k + n_k)$$

- Define $\Delta = \sum_k \frac{(p_k - \hat{p}_k)^2}{\hat{p}_k} + \frac{(n_k - \hat{n}_k)^2}{\hat{n}_k}$, Δ is of χ^2 distribution with $(n + p - 1)$ degree of freedom.
- For example, with 3 degree of freedom, $\Delta \leq 7.82$ encourages the pruning with 5% **level of significance**. 不夠大->破了

Rule Post-Pruning

J48 v8

- Used by C4.5rules [Quinlan, 1993].

tree 沒有避免 overfitting

- Convert the decision tree into rules (one rule per path).
 - Prune each rule by removing any preconditions that improves its accuracy (by testing set). 砍砍看 precondition accuracy 上升, 則take
 - Sort the the pruned rules by their accuracy, and consider them in this sequence when classifying instances.
- For example,
rule $(Patron = Full) \wedge (Hungry = No) \Rightarrow (WillWait = False)$.
 - Rule post-pruning considers removing $(Patron = Full)$ and $(Hungry = NO)$ in this example.

Ensemble

- Use multiple **weak classifiers** to prevent from overfitting.
- Embedding
- Bagging
- Boosting



Embedding

≈ Kernel method

- Random forest: randomly select k attributes to create weak classifiers.

- Forest is ensemble of several decision trees

total: # of attribute
對於k值的選擇
sqrt() or log()



– classification is
$$P(c|\mathbf{v}) = \frac{1}{T} \sum_{t=1}^T P_t(c|\mathbf{v})$$

[Amit & Geman 97]
[Breiman 01]
[Lepetit et al. 06]

Bagging

- Sampling with replacement from the dataset to form new datasets.

Original Data	1	2	3	4	5	6	7	8	9	10
Bagging (Round 1)	7	8	10	8	2	5	10	10	5	9
Bagging (Round 2)	1	4	9	1	2	3	2	7	3	2
Bagging (Round 3)	1	8	5	10	5	5	9	6	3	7

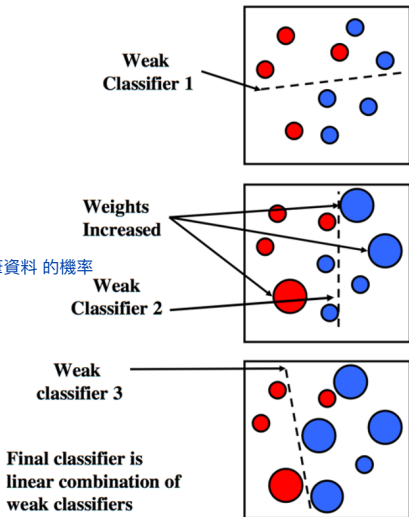
- Build classifier on each bootstrap sample (supposedly n items).
- Probability $(1 - 1/n)^n$ of not being selected. n 個通通沒被選到
- When n is large, it is about $1/e \simeq 37\%$
- About 37% of noise (if any) not being selected.

好處：有可能完全沒有noise

Boosting

if preprocessing is good, boosting
if not, bagging is better.

- An iterative procedure to adaptively change distribution of training data by focusing more on previously misclassified records.
- Initially, all n items are assigned equal weights.
- Unlike bagging, weights vary at the end of boosting round.



AdaBoost

- Weak classifiers: C_i
- Error rates:

$$\epsilon_i = \frac{1}{N} \sum_{j=1}^N w_j \cdot \delta[C_i(x_j) \neq y_j]$$

- Importance of a classifier:

錯誤率越小 越重要

$$\alpha_i = \frac{1}{2} \ln \frac{1 - \epsilon_i}{\epsilon_i}$$

- Weight update (c normalization factor):

$$w_j \leftarrow c \cdot w_j \begin{cases} e^{-\alpha_i} & C_i(x_j) = y_j \\ e^{\alpha_i} & C_i(x_j) \neq y_j \end{cases}$$

對錯 縮小放大

AdaBoost

- The equations of the previous slide are such to minimize the total error. We omit the derivations here.
- Initially, $w_j = \frac{1}{N}$.
- Any intermediate round yields error rate higher than 0.5, weights are reverted back to $\frac{1}{N}$.
- Classification:

$$C^*(x) = \operatorname{argmax}_y \sum_j \alpha_j \cdot \delta[C_j(x) = y]$$

Summary

- **Decision tree** learning using information gain.
- Learning performance = prediction accuracy measured on test set.
- **Cross-validation** combats **overfitting**.
- **Bayesian learning** is reasonable, but computationally expensive. A common simplification is the **MAP** learning.
- **MAP** learning reduces to finding the **ML** hypothesis when assuming all hypotheses are equally probable.
- Prior that penalizes complexity combat **overfitting** and results in **MDL**.
- **Bayesian networks** provide a natural representation for (causally induced) conditional independence.
- Topology + CPTs = compact representation of joint distribution.
- Bayesian networks are generally easy to construct with human knowledge or other **machine learning** techniques.